

Appendix POS — The Principle of Sequence (v1.6.1)

APPENDIX POS: THE PRINCIPLE OF SEQUENCE — READINESS-ORDERED COMMITMENT

The Theory of the Forbidden Action

Copyright © 2026 Petar Nikolov. Licensed under CC BY 4.0.
Appendix for U-Theory / U-Model — v.28 appendix series | Aligned with SSS / LGP / GSI-RTD / TE / TAA
Status: L2 methodology — a **commitment-control discipline** layered on the SSS measurement engine
Version: 1.6.1
Epistemic Level: B2 / L2 — a disciplined commitment heuristic; the claim that readiness-gated commitment beats ungated alternatives at equal cost is a **pre-registered hypothesis** (§8), not yet a validated result
Author: Petar Nikolov (ORCID 0009-0001-8669-2276)
Prerequisites: [APPENDIX_SSS](#) (the score), [APPENDIX_LGP](#) (the cycle), [APPENDIX_GSI-RTD](#) (the runtime), [APPENDIX_TE](#) (the umbrella)
v26 Invariant: Form ↔ Time · Position ↔ Space · Action ↔ Energy

Leading thought — P. Nikolov · tagged B0 / VISION (renders no result; a motivation, not a finding).

Take an athlete out of shape and place him at the start of the final. He will not win — not because he lacks will, but because the body was never built for the moment. First the form, then the lane, then the gun. The mistake we make in life is almost never that we acted too weakly; it is that we committed too early — from a body not yet built, or a place that could never carry the blow. $U = \sqrt[3]{(F \cdot P \cdot A)}$ tells you what to grow. This appendix tells you when a decisive step may be taken, and the one commitment you must refuse until the base is ready.

0. What POS is, what it is NOT, and what is new

The keystone $U = \sqrt[3]{(F \cdot P \cdot A)}$ answers **what to maximize**: the *product* of the three pillars. Accumulated as $\mathcal{M} = \int U \, dt$, that product is the system's **cumulative stability** — the area under its U -trajectory. (Its further reading as *meaning* belongs to [APPENDIX_MMT](#) / [APPENDIX_DIM](#) and inherits their L3 status — §6.) But the keystone defines a **score over states**; it does not define a **policy for committing to actions in time**. The **Principle of Sequence (POS)** supplies exactly that missing piece — a *commitment gate*.

The one-sentence core, stated at its most defensible:

POS does not dictate the order in which a system thinks or generates candidate solutions. POS dictates the conditions under which an already-generated decision may pass from reversible exploration to irreversible commitment: Form-readiness and Position-readiness are verified — in that order — before the decisive action is authorized.

Three distinct orders coexist in the corpus, and POS is the third — it replaces neither of the first two:

Layer	Order	Owner
Ontology — what exists	$F \parallel P \parallel A$ (co-equal, simultaneous)	the core invariant
Discovery / candidate generation — how to search	$A \rightarrow F \mid A \rightarrow P \mid F, A$ (AD-RTD: <i>what must be done?</i> $\rightarrow$ with <i>what?</i> $\rightarrow$ <i>where/when?</i>)	GSI-RTD §1.1
Commitment — when to execute irreversibly	<code>check F $\rightarrow$ check P $\rightarrow$ authorize execute(a)</code>	POS (this appendix)

- **SSS** measures $F, P, A \rightarrow$ POS **consumes** those measurements as gate inputs.
- **GSI-RTD (AD-RTD)** generates and triages candidates Action-first — the right order for discovery, because Action defines the problem-solving intent (GSI-RTD Prop. 1.1). $\rightarrow$ POS adds the **commitment check** AD-RTD does not itself supply: a candidate may be *found* Action-first, but may only be *executed* readiness-first.
- **LGP** diagnoses F, P, A **in parallel** (LGP-1...3 are orthogonal by design) and *selects* a solution at LGP-7. $\rightarrow$ POS is the **gate between LGP-7 (selection) and LGP-8/9 (plan & resource commitment)**: the selected solution passes to execution only if the readiness check clears.

POS is: a readiness-ordered **commitment discipline** — a rule for *when* a decisive, hard-to-reverse action may be executed, so that U stays high across time and $\mathcal{M}$ keeps accumulating.

POS is NOT: a new pillar, a new aggregator, a thinking-order mandate, a serialization of LGP's parallel diagnostics, or a replacement for AD-RTD's Action-first discovery. It does not license any action that a values-firewall forbids (§7). It is a *default gate under uncertainty*, with an explicit forced-move exception (§8).

What is new: an explicit, teachable **commitment gate with a named failure mode** — the *Forbidden Action* (the unprepared irreversible commitment) — an action taxonomy that dissolves the "you must act to build Form" circularity (§2.5), and a pre-registered test of whether the gate pays (§8).

1. The Principle

■ THE PRINCIPLE OF SEQUENCE

Never authorize a decisive, hard-to-reverse action (a^{commit}) until you have verified, in order:

(1) your *Form* is fit for this action, and (2) your *Position* can carry it.

The unprepared commitment is the *Forbidden Action*: good Form burns up in a bad Position, U collapses toward zero, and the growth of $\mathcal{M}$ collapses with it.

Equivalently, in the currencies of the invariant: **the decisive commitment is placed last** — not because energy is destroyed (it is converted, not lost), but because a large commitment **spends limited resources, forecloses future options, and changes the system's state in ways that are costly or impossible to undo**. The sequence protects the base that every future U is multiplied against. (Small, reversible actions are *not* gated — see §2.5.)

2. The three gates — the commitment review

POS runs as a **cyclical readiness review over a specific candidate action**. Facing a decisive step — in a project, a company, a body, a nation — you pass three gates *in order* **before committing**. The gates order the *verification*, not the thinking: diagnosis of all three pillars may (and in LGP does) run in parallel; the *authorization* is sequential.

Gate 1 — FORM (verify first, always)

Before any decisive commitment, verify **Form** against *this specific action*:

- Is my Form **fit for the goal this action serves** — not in the abstract, but for *this* objective?
- Can this Form **survive the Position** it will be exposed to — an aggressive environment, hostile conditions, a demanding market?
- Can this Form **support this action's demands** without tearing?

If any answer is *no*, the commitment waits. You repair the Form first — using preparatory actions (permitted freely only when low- κ ; §2.5).

Gate 2 — POSITION (earn it, with Form in hand)

Only with a Form fit for the action do you verify — and work to improve — **Position**. Position depends heavily on Form; you cannot hold a place your Form cannot defend.

The athlete. Take a runner in poor condition and hand him lane one of the final — the best position on the track. Can he win? No. First he must **build his form** — train, add the muscle, sharpen the engine — and *then* compete for the place. Position is the competition; you enter it only once the Form is built. To seize a position your Form cannot defend is to lose both.

Gate 3 — ACTION (the decisive step, on a prepared base)

Only with a ready Form and a secured Position is the decisive action **authorized**. And the door opens wide: some actions you planned from the start; **others become available only because you spent so long on Form and Position** — opportunities that simply do not exist for the unprepared.

"More sweat in training, less blood in battle" (attributed to A. V. Suvorov). The decisive action is the start of the battle — the moment you commit. It achieves its result **only** on a prepared base. The blood you save in the fight is the sweat you paid in the Form and the Position beforehand.

2.4 Readiness thresholds — action-, domain-, and stakes-specific

The gates need a *when*. POS defines the thresholds as **functions, not constants**:

$\theta_F(a, d, s)$ and $\theta_P(a, d, s)$ — the minimal Form- and Position-readiness for **action a , in domain d , at stakes s** — compared against a **lower confidence bound** on the measured pillar (LCB, not an optimistic point estimate; uncertainty-aware per the TRA/TRC GSM discipline).

- Higher stakes / lower failure tolerance $\rightarrow$ higher θ (SSS's own stakes-scaling: ordinary contexts near the $\varphi^{-1} \approx 0.618$ band; critical infrastructure ≥ 0.75 — **SSS reference values, cited for scale only; no independent decision authority in POS; $\theta(a, d, s)$ calibration governs**).
- For any *irreversible* verdict, GSI-RTD's SSS-Guard additionally requires the SSS verdict to be cross-checked against an independent domain metric — POS inherits that requirement unchanged.
- **Illustrative placeholders with ZERO decision authority** (retained only to make the mechanism concrete): "Gate 3 at $F \gtrsim 0.6$ and $P \gtrsim 0.5$ " — un-calibrated, asymmetry unjustified, superseded by $\theta(a, d, s)$ the moment POS-P1 calibration exists.

2.5 The action taxonomy — dissolving the circularity

"To repair Form you must act — but action is forbidden until Form is ready" would be a genuine paradox **if POS gated all action**. It does not. POS distinguishes the pillar-score from concrete actions, and gates only one class:

Sym bol	Meaning	Gated by POS?
A_t	the Action pillar score — current operational capability	not an action at all; a measurement
a^p_{rob}	a small, cheap, reversible exploratory action (a test, a pilot, a conversation, a paper trade)	No — capped exposure, always permitted
a^p_{rep}	an action that builds Form or Position (training, hiring, refactoring, coalition-building)	Only if high-κ — low- κ prep flows freely (this <i>is</i> the sequencing work); a large, hard-to-reverse "preparation" (mass hiring, capital build, irreversible restructuring) is a commitment in disguise and must clear the gate
a^c_{ommit}	a decisive, hard-to-reverse commitment (the race, the scale-up, the reform, the launch)	YES — the POS gate applies here and only here
a^e_{mg}	a forced move in a genuine emergency (§8)	Exempt, but must be logged as forced

Each action carries an **irreversibility degree** $\kappa(a) \in [0, 1]$ (0 = fully reversible probe; 1 = practically irreversible commitment). The gate, compactly:

a is authorized iff the values-firewall passes (§7), AND at least one of:

- (i) $\kappa(a) \leq \kappa_{\text{probe}}$ (it is a reversible probe/preparation — proceed with capped exposure);
- (ii) $\text{LCB}(F|a) \geq \theta_F(a, d, s)$ and $\text{LCB}(P|a) \geq \theta_P(a, d, s)$ (the base is ready — commit);
- (iii) **expected loss of inaction exceeds expected loss of acting, and the move is logged** a^{emg} .

Unprepared + irreversible + not an emergency = the Forbidden Action. Everything else flows.

No label-gaming. Gate applicability keys on $\kappa(a)$, exposure, and stakes — *not* on what the action is called. Naming a high- κ commitment "preparation" does not exempt it; and the values-firewall (§7) dominates at every κ .

2.6 Formal core — the gate as a function

The gate is not only prose; it is a computable predicate. (*Plain-symbol notation, not LaTeX; $\vartheta_F(a)$ is shorthand for the action·domain·stakes-conditioned $\vartheta_F(a, d, s)$ of §2.4.*)

Readiness ratio — how far a candidate action clears **both** thresholds, measured on the *lower confidence bound* of each pillar (never an optimistic point estimate):

$$R(a, t) = \min(\text{LCB}(F_t | a) / \theta_F(a) , \text{LCB}(P_t | a) / \theta_P(a))$$

$$R(a, t) \geq 1 \quad \Leftrightarrow \quad \text{both Form- and Position-readiness clear their thresholds at the chosen confidence standard.}$$

2.6.1 Readiness margin and the active bottleneck

$R(a, t)$ says *whether* both conditions clear; the **margin** says *by how much*:

$$\begin{aligned} m_F(a, t) &= \text{LCB}(F_t | a) / \theta_F(a, d, s) - 1 \\ m_P(a, t) &= \text{LCB}(P_t | a) / \theta_P(a, d, s) - 1 \\ m_R(a, t) &= \min(m_F , m_P) = R(a, t) - 1 \end{aligned}$$

$m_R > 0$: both clear with positive reserve · $m_R = 0$: exactly on the boundary · $m_R < 0$: at least one unmet. The **active readiness bottleneck** $b(a, t) = \text{argmin}_{i \in \{F, P\}} m_i(a, t)$ names the limiting pillar (F , P , or a tie). It authorizes nothing; it points preparatory effort at the pillar with the highest immediate readiness value.

2.6.2 Confidence standard

POS reads each pillar at a *lower confidence bound*, never an optimistic point estimate. For $x \in \{F, P\}$:

$$\begin{aligned} \text{LCB}_{\{1-\alpha\}}(X_t | a) &= Q_{\alpha}(X_t | a, D_t) && \text{(lower } \alpha\text{-quantile of the posterior / sampling dist.)} \\ &\approx \text{mean}[X_t(a)] - z_{\{1-\alpha\}} \cdot \text{SE}[X_t(a)] && \text{(approximately-Gaussian estimator)} \end{aligned}$$

A one-sided 95% bound uses $\alpha = 0.05$. The requirement scales with stakes: $\alpha = \alpha(s)$, $d\alpha/ds \leq 0$ — higher stakes demand a more conservative bound. Estimator, uncertainty model, and α are **declared before the gate is evaluated**; POS may not pick a friendlier uncertainty model after seeing whether the action passes.

2.6.3 Risk-adjusted readiness thresholds (a calibration candidate)

θ_F , θ_P must ultimately be calibrated (POS-P1). A candidate family to *test*:

$$\theta_i(a,d,s) = \text{clip}_{[0,1]}(\theta_i^o(d) + \beta_i \cdot \kappa(a) + \gamma_i \cdot e(a) + \eta_i \cdot s) , \quad i \in \{F,P\}$$

with $\theta_i^o(d)$ the domain baseline, $\kappa(a)$ irreversibility, $e(a) \in [0,1]$ normalized exposure, $s \in [0,1]$ normalized stakes, $\beta_i, \gamma_i, \eta_i \geq 0$ calibrated coefficients. Intended monotonicity $\partial \theta_i / \partial \kappa \geq 0$, $\partial \theta_i / \partial e \geq 0$, $\partial \theta_i / \partial s \geq 0$ — the more irreversible, exposed, or consequential the action, the stronger the readiness evidence required. A **calibration candidate, not a validated law**; POS-P1 tests it against simpler and non-linear rivals.

Degree of irreversibility $\kappa(a) \in [0,1]$: $\kappa=0$ a fully reversible probe; $\kappa=1$ a practically irreversible commitment; intermediate = partial reversibility / recovery cost / limited option preservation.

The POS authorization gate $G_POS(a,t) \in \{0,1\}$ — authorize (1) **iff** the values-firewall passes **and** at least one readiness condition holds:

$$\begin{aligned} G_POS(a,t) = & 1[\text{Firewall}(a) = \text{PASS}] \\ & \cdot 1[\kappa(a) \leq \kappa_probe] \quad (\text{reversible} \rightarrow \text{capped exposure}) \\ & \vee R(a,t) \geq 1 \quad (\text{prepared} \rightarrow \text{commit}) \\ & \vee (E[L(\text{inaction})] > E[L(a)] \wedge a = a^{emg}) \quad (\text{forced move} \rightarrow \text{logged}) \end{aligned}$$

Condition	POS verdict
Unprepared + reversible probe ($\kappa \leq \kappa_probe$)	permitted, capped exposure
Prepared + decisive action ($R \geq 1$)	permitted
Unprepared + decisive action (κ high, $R < 1$)	readiness-forbidden — the Forbidden Action
Genuine emergency (a^{emg})	forced move — mandatory logging
Values-forbidden action	always prohibited

The **values-firewall dominates every readiness result**: no $R(a,t)$ or $\kappa(a)$, however favorable, authorizes an action barred by ethics, law, or a domain safety constraint (§7).

2.6.4 Severity of a readiness violation

The binary gate says authorized / not; it does not grade *how bad* an attempted violation is. For audit and prioritization:

$$V_POS(a,t) = s(a) \cdot \kappa(a) \cdot [1 - R(a,t)]_+ , \quad [x]_+ = \max(x, 0)$$

Properties: $R \geq 1 \Rightarrow V_POS = 0$; $\kappa \rightarrow 0 \Rightarrow V_POS \rightarrow 0$; and $R \ll 1, \kappa \rightarrow 1, s \rightarrow 1 \Rightarrow V_POS$ large — the worst violation is a high-stakes, highly irreversible commitment far below the boundary.

V_POS is an **audit measure, not a second gate**, and a values-firewall failure is *never* folded into it (values-forbidden actions stay categorically prohibited regardless of V_POS).

2.6.5 Uncertainty-adjusted emergency override

The emergency branch must not fire merely because the *point estimate* of inaction loss exceeds that of action. Require a **robust** margin on confidence bounds:

$$\begin{aligned} M_emg(a,t) &= \text{LCB}[L(\text{inaction})] - \text{UCB}[L(a)] && \text{emergency iff } M_emg > 0 \\ ER(a,t) &= (\text{LCB}[L(\text{inaction})] + \varepsilon) / (\text{UCB}[L(a)] + \varepsilon) , && \text{robust iff } ER > 1 \quad (\varepsilon > 0) \end{aligned}$$

The gate's emergency branch tightens to $a = a^{emg} \wedge M_emg(a,t) > 0 \wedge \text{Log}(a,t) = \text{COMPLETE}$. This stops an optimistic estimate of action success from manufacturing an emergency exemption.

Where time forbids reliable bounds, the override is recorded as taken under explicit, unresolved uncertainty.

3. The Forbidden Action

The **Forbidden Action** is the sharp, teachable failure mode POS names: **a decisive, hard-to-reverse commitment (a^{commit}) taken before Form and Position are ready for it.**

Why it earns the name — stated honestly:

1. **Non-compensation at the limit.** In $U = \sqrt[3]{(F \cdot P \cdot A)}$, if a pillar is *near zero*, no surplus elsewhere rescues the product — the zero-limit is where non-compensation is exact. **In the interior, the geometric mean permits partial trade** ($F=0.5, P=1, A=1$ still gives $U \approx 0.794$) — so the gate is **not a theorem derived from U** ; it is a **normative engineering constraint** (per the TRA hardening: away from zero, U is a penalized trade, not a hard limit). The gate must therefore *earn* its keep empirically — that is exactly POS-P1 (§8).
2. **The base burns.** A large commitment from an unready base spends its resources *on* the base — consuming it — rather than *from* a surplus. What makes the loss stick is not "energy irreversibility" in the thermodynamic-slogan sense (energy converts; what is lost is *available* work, options, and state), but **commitment irreversibility**: $\kappa(a^{\text{commit}}) \rightarrow 1$ means the resources, options, and standing consumed do not come back on any useful timescale (§3.1).
3. **Accumulation collapses in rate.** $\mathcal{M} = \int U \, dt$ is monotone non-decreasing for $U \geq 0$ — it cannot literally reverse. What the Forbidden Action destroys is its **growth**: a burned base drops U and holds it down, so $d\mathcal{M}/dt = U$ collapses, and $\mathcal{M}$ **falls permanently behind the trajectory the prepared alternative would have accumulated**. (Where an explicit stock that *can* decline is needed, model it separately as $\mathcal{M}_{\text{net}} = \mathcal{M}_0 + \int [U(\tau) - L(\tau)] d\tau$ with a declared loss term L ; the un-normalized $\mathcal{M}$ itself only slows or stalls.)

The mirror of the prohibition is its promise: **an action that looks impossible or "forbidden" from your present base is very often a readiness problem, not a permanent ban.** Sequence the Form, then the Position, and the step that was unreachable becomes a natural, low-blood move. The Principle does not teach you to force forbidden doors; it teaches you to build until the right doors open on their own.

3.1 Mechanisms of base burning — three scenarios

The claim that a premature commitment *degrades* the base (not merely fails to advance) is concrete, and the mechanism is the same across substrates: **the commitment consumes limited resources and forecloses options, and at $\kappa \rightarrow 1$ that consumption does not reverse.**

- **Physical (the athlete, continued).** Racing un-recovered does not merely lose the race; the unprepared body, forced to perform, **tears** — the commitment injures the Form, leaving it *below* its pre-race level. The base is not just unused; it is damaged.
- **Business (the premature scale-up).** A start-up that scales — aggressive spend, hiring, expansion — before product-market fit (Form) and a defensible niche (Position) burns its capital *and* its reputation. The failed launch poisons the brand, so even a later, genuinely good product cannot get a hearing: credibility drops *below* baseline and the spent capital does not return. The base is burned twice.
- **Institutional / political (the premature reform).** A government that launches a major reform before institutional capacity (Form) and legitimacy / coalition (Position) are ready does not

merely fail; the botched reform **discredits the institutions** (Form degraded) and **burns political capital** (Position lost), leaving the goal *further away* than before the attempt.

In each case the signature is identical: **U** does not merely dip — it **drops and stays down**, so the growth of **M** collapses and the system falls behind the accumulation the sequenced path would have delivered. (Causal attribution is exactly what POS-P2 must establish against a sequenced comparator — observation of damage alone does not prove the mechanism; §8.)

3.2 The base-burning metric — and why it needs a causal comparator

The "burn" is measurable. Define the immediate base-damage after an action as the weighted, one-sided drop in the Form/Position base:

$$B(a,t) = w_F \cdot [F_t - F_{t+1}]_+ + w_P \cdot [P_t - P_{t+1}]_+, \quad [x]_+ = \max(x, 0)$$

with domain-calibrated weights $w_F, w_P \geq 0$ **declared before analysis**. But $B(a,t) > 0$ **does not by itself prove the action caused the damage** (regression to the mean, confounders). POS-P2 therefore compares two *policies* under matched initial conditions, randomized assignment, or a causally adjusted design:

$$\Delta B = E[B \mid \pi_{\text{premature}}] - E[B \mid \pi_{\text{sequenced}}]$$

$\Delta B > 0$ supports the base-burning mechanism (premature commitment damages the base more than the sequenced policy); $\Delta B \leq 0$ weakens or falsifies it, under the pre-registered equivalence / uncertainty criteria (§8, POS-P2).

3.3 Base dynamics — separating construction from destruction

The one-step burn records a decline but does not split gains from damages. Model each pillar as gain minus damage:

$$\begin{aligned} F_{t+1} &= \text{clip}_{[0,1]}(F_t + g_F(a,t) - d_F(a,t)) \\ P_{t+1} &= \text{clip}_{[0,1]}(P_t + g_P(a,t) - d_P(a,t)) \end{aligned} \quad g_*, d_* \geq 0$$

Immediate burn $B_1 = w_F \cdot d_F + w_P \cdot d_P$. To separate temporary disruption from *persistent* damage, use the discounted H -horizon burn:

$$B_H(a,t) = \sum_{h=1..H} \gamma^{h-1} (w_F \cdot [F_t - F_{t+h}]_+ + w_P \cdot [P_t - P_{t+h}]_+), \quad 0 < \gamma \leq 1$$

A dip that fully recovers yields a smaller B_H than a persistent decline. POS-P2 preregisters H, γ, w_F, w_P , and the minimum practically-important B_H .

3.4 Counterfactual accumulation gap

"Falls behind the prepared trajectory" made precise — the gap between the sequenced and premature policies:

$$\Delta M_T = \int_{[0,T]} (U_{\text{seq}}(t) - U_{\text{prem}}(t)) dt \quad (\text{normalized: } \Delta M_T / T)$$

$\Delta M_T > 0$: the sequenced policy accumulated more stability • $= 0$: no advantage • < 0 : the ungated policy accumulated more. Because both trajectories come from the *same* internal pillar

model, ΔM_T is a **secondary internal outcome**; POS-P1 keeps external domain outcomes as its primary evidence.

4. Why F-then-P is the default verification order — and why this does not touch AD-RTD

Both readiness conditions must hold before a commitment, so the *conjunction* $LCB(F|a) \geq \theta_F \wedge LCB(P|a) \geq \theta_P$ is itself **order-independent**. POS nonetheless verifies **F first** as its canonical **default** — motivated by dependency and check-economy, *not* claimed as a pre-proven structural necessity (whether F-first strictly beats P-first is itself part of POS-P1, §8):

- **Dependency & check-economy. Position is earned through Form** (the athlete): a Position your Form cannot hold is lost the moment it is tested, so a P-check can be invalidated by a later F-failure. Checking F first lets an early F-failure abort **cheaply**, before the often-costlier P-assessment; and P-readiness is frequently ill-defined until a minimal Form exists. Order therefore bites precisely when the checks differ in cost, when P-validity depends on confirmed F, or when early rejection saves resources. ($P \rightarrow F \rightarrow A$ collapses back into "fix Form first" — after a wasted move.)
- **Option value.** The decisive commitment is verified last and executed last because it is the **least reversible** move — the one that forecloses the most options. Reversible work (probes, preparation) keeps options open; the commitment spends them. You spend option-value last.
- **Integral, not instant.** POS optimizes the accumulation $\mathcal{M}$, not a single moment's U . The gate that protects the base maximizes the *area under the curve* across the whole life of the system — not one brilliant, base-burning moment.
- **No conflict with Action-first discovery.** AD-RTD's $A \rightarrow F \rightarrow P$ answers "how do I efficiently find what to do?" — Action defines intent, so discovery correctly starts there (GSI-RTD Prop. 1.1: reversing the discovery order "generates structure without purpose"). POS answers a different question: "when may the found action be executed irreversibly?" **Discover Action-first; commit readiness-first.** The two orders operate on different objects (candidate-space vs. commitment) and compose cleanly — §5.
- **Coupling honesty.** In some systems (notably political ones) available *Forms* are constrained by *Position*, so F and P must co-evolve. POS handles this precisely: **F-first is the canonical default verification order; not an unconditional universal requirement.** A domain adapter MAY select P-first when a pre-registered verification-cost or dependency analysis establishes $C_{\{P \rightarrow F\}} < C_{\{F \rightarrow P\}}$ (§4.1), or when Position must be resolved before Form can be validly estimated; the selected order, its cost/dependency assumptions, and the adapter version are recorded in the decision certificate. This changes only the *verification schedule* — the **commit-gate is unconditional** and the final readiness predicate is order-independent:

$$LCB(F|a) \geq \theta_F \quad \wedge \quad LCB(P|a) \geq \theta_P$$

The *building* of F and P may interleave freely via a^{prep} (§2.5, §8); order-of-verification and order-of-build are distinct.

4.1 The economics of checking Form first

Since $F\text{-ready} \wedge P\text{-ready}$ is order-independent, the value of checking Form first comes from cost, information, or dependency. Let c_F, c_P be the verification costs and p_F, p_P the pass-probabilities. With early-exit on the first failure:

$$C_{\{F \rightarrow P\}} = c_F + p_F \cdot c_P$$

$$C_{\{P \rightarrow F\}} = c_P + p_P \cdot c_F$$

F-first is economically preferred when $C_{\{F \rightarrow P\}} \leq C_{\{P \rightarrow F\}}$, i.e. $c_F \cdot (1 - p_P) \leq c_P \cdot (1 - p_F)$ — a cheap gate that often rejects should be checked early; an expensive gate deferred until cheaper necessary conditions pass. When a Position estimate becomes invalid after a Form failure, add a dependency term:

$$C^*_{\{P \rightarrow F\}} = C_{\{P \rightarrow F\}} + c_{\text{invalid}} \cdot q_{\{P | \neg F\}}$$

So **F-first's superiority is an empirical proposition** about costs, rejection probabilities, and dependency — not a consequence of the conjunction alone.

4.2 Commitment burden and option-foreclosure

Irreversibility is not only physical recovery cost; a commitment can destroy valuable future choices. Define:

$$\begin{aligned} \Omega(a, t) &= V_{\text{free}}(t) - V_{\text{after}, a}(t) \\ C_{\text{commit}}(a, t) &= \kappa(a) \cdot (c(a, t) + \lambda_{\Omega} \cdot \Omega(a, t)), \quad \lambda_{\Omega} \geq 0 \end{aligned}$$

where V_{free} is the value of the best future policy *before* commitment and $V_{\text{after}, a}$ *after* committing a . A large immediate payoff can still carry a high burden when it prunes superior future branches. This formalizes the rule: **reversible exploration preserves the policy tree; irreversible commitment prunes it.** Ω is model-dependent — its estimation method and sensitivity range are declared before use.

5. Reconciliation & integration — where POS sits in the suite

The full pipeline, with each protocol doing what it is for:

- | |
|---|
| 1. DISCOVERY / CANDIDATE GENERATION
GSI-RTD (AD-RTD): $A \rightarrow F A \rightarrow P F, A$
"what must be done? → with what? → where?" |
| 2. PARALLEL TRIADIC ASSESSMENT
TAA / LGP-1...3: evaluate $F \parallel P \parallel A$
(orthogonal diagnostics – never serialized) |
| 3. RANKING & SELECTION
LGP-4...7: impact → leverage → synth → select |
| 4. ★ POS COMMITMENT GATE ★
check F → check P → authorize execute(a)
probes/prep flow free; a^commit must clear;
SSS-Guard cross-check on irreversible verdicts |
| 5. EXECUTION & CONTROL
LGP-8/9 (plan, resources) → LGP-10 monitoring
(abort discipline) → LGP-12 / SSS final audit |

Appendix	Its role	What POS adds to it
SSS	measures U , δ , SI (+ uncertainty)	the commitment gate consuming its LCBs; a^{commit} may not execute until F and P clear $\theta(a, d, s)$
GSI-RTD	AD-RTD candidate generation ($A \rightarrow F \rightarrow P$), the Triadic Scheduler (fed by π^* weakest-pillar triage), SSS-Guard	the commitment check AD-RTD does not itself supply ; POS inherits SSS-Guard for irreversible verdicts
LGP	12-step cycle; parallel F/P/A diagnostics (LGP-1...3); selection (LGP-7); abort (LGP-10)	the gate between LGP-7 and LGP-8/9 — selection may pass to plan/resources only through the readiness check
TAA	the F/P/A/ Σ agent jury	the Σ agent's commitment veto — Σ withholds authorization of a^{commit} until the F and P agents report ready
TE	the engineering umbrella	POS as a cross-cutting commitment discipline for every domain instance (TRA, TRB, TRC, TSE)

Canonical execution order (unchanged): GSI-RTD $\rightarrow$ TAA $\rightarrow$ LGP-12 $\rightarrow$ SSS. POS does not reorder it; POS inserts one checkpoint inside it.

6. Universality — a working hypothesis

The gate is claimed to be **substrate-independent** — *a working hypothesis, tested by POS-P1 (§8), not an established fact* — applying wherever $U = \sqrt[3]{(F \cdot P \cdot A)}$ applies:

- **Personal life.** Meet a decisive step by verifying your **Form** first (health, skill, character, resources of the self), then your **Position** (environment, relationships, standing, timing), and only then **commit**. Most personal disasters are Forbidden Actions: a decisive move made from an unready self or an impossible position. Probes and preparation are always allowed — that is how the base gets built.
- **Public life and governance.** The Universal Model's civilizational aims — **higher productivity, better service to citizens, greater happiness, lower mortality** — are served the same way: build institutional **Form** (competent, honest structures), secure **Position** (stable environment, legitimacy, resources), and only then launch decisive **Action** (reforms, programs, commitments). A reform committed from an unready institutional base burns the base and sets the goal back.

L3 interpretive note (non-load-bearing). In the MMT/DIM reading, a reasoning system that sequences its commitments turns raw capability into durable, accumulating meaning — the greatest attainable density of matter and meaning — rather than a single bright flare. This is an interpretation layer (`APPENDIX_MMT` is L3-speculative); nothing in POS's operational content or falsifiers rests on it.

7. Two meanings of "forbidden" — readiness vs. values

The word *forbidden* in this appendix has a **precise, single meaning**: an action is *readiness-forbidden* when Form or Position is not yet ready to carry it. Such an action is not banned forever — it becomes available once the base is built (§3).

This must never be confused with **values-forbidden** action — actions barred by ethics, law, and the corpus's dual-use firewall (`APPENDIX_WAR`, the TRC/TRA firewalls). **POS never licenses a values-**

forbidden action. Sequencing Form and Position can unlock a *hard* action; it can never make a *wrong* one legitimate. The values-firewall **dominates the gate**: it is checked first (§2.5), and a values-forbidden action stays forbidden at every level of readiness.

8. Falsifiability and honest scope

POS is a **disciplined heuristic**, held to the corpus's own evidence standard.

What POS claims: that gating decisive commitments on F-then-P readiness (with probes and preparation ungated) yields better long-run outcomes than ungated alternatives, at equal total budget.

What POS does NOT claim: that it guarantees success; that action is never urgent; that F, P, A are cleanly separable in every system; that the gate follows as a theorem from **U** (§3.1); or that it replaces domain expertise.

POS-P1 (the gate pays) — protocol sketch (full preregistration required before any run: endpoint, N, power, minimum effect $\Delta_{\min}$, analysis plan, public registration record):

- **Baselines (mandatory):** (i) **AD-RTD alone** (Action-first discovery, no commitment gate) — *the incumbent GSI canon and the baseline of record*; (ii) weakest-pillar π^* -greedy; (iii) expert/domain policy; (iv) explore-then-commit; (v) random order; (vi) naive action-first; (vii) **thresholds-without-order** and (viii) **order-without-gate** — the ablations that isolate *which component* (order, gate, or mere caution) carries any effect.
- **Outcomes:** primarily **external** domain metrics (task success, failure/damage rate, cost, time, survival, cumulative regret) — *not* $\mathcal{M}$ computed from the same F/P/A estimates the gate consumes (that would be partially self-confirming).
- **Equal budget:** total = time + resources + information + compute + **risk exposure**.
- **Verdicts:** **PASS** iff $LCB_{95}(\text{effect vs. best baseline}) > \Delta_{\min}$; **KILL** iff $UCB_{95} < 0$ (the gate *harms*); everything between is **INCONCLUSIVE** — neither victory nor death, and reported as such.
- **Heterogeneity:** effects estimated per-domain (hierarchical), because "universal" is exactly what must be tested, not assumed.

POS-P2 (the base burns) — causal, not observational: compare $\Delta B = E[\text{base-damage} \mid \text{premature } \pi] - E[\text{base-damage} \mid \text{sequenced } \pi]$ under matched initial conditions or a causally adjusted design; damage observed after premature commitment alone proves nothing (regression to the mean, confounders). **KILL:** $\Delta B \leq 0 \rightarrow$ the "burns" mechanic is demoted to metaphor and §3.1 is re-labeled illustrative.

Load-bearing caveat — coupling. In some systems Position can *only* be improved **by** acting; POS already routes this through **a^{prep}** (ungated). Where even large commitments are the only way to move P, POS is a **default gate under uncertainty**, not an absolute prohibition — and says so.

Forced-move exception. When *not* acting is itself the larger, irreversible loss (a genuine emergency), the disciplined move is to act on the best available base — logged as **a^{emg}**, a forced move, not a Forbidden Action. The log requirement prevents the exception from quietly swallowing the rule (every **a^{emg}** is auditable at LGP-12).

Apparent counterexamples (why they do not falsify POS). Celebrated "act-first" strategies — "*fake it till you make it*," the scrappy launch, the bluff that buys time — are almost never unprepared *commitments*. Examined against the taxonomy (§2.5) they are **a^{probe}** / **a^{prep}**: small, cheap, **reversible** moves taken to *improve* Position or Form, on a base that survives their failure. POS endorses them. POS is falsified only if **unprepared a^{commit}** — decisive, irreversible,

from a genuinely unready base — **systematically outperforms** the gated alternative at equal budget (POS-P1). The reclassification discipline cuts both ways: an action may be called a probe **only before** its outcome is known — no post-hoc relabeling of failed commitments as "probes."

POS-P3 (advanced research formalization) — the policy-level control objective.

RESEARCH OBJECTIVE ONLY — NOT FOR OPERATIONAL AUTHORIZATION. The policy-level objective below is a mathematical research specification, not a deployment rule. It must not authorize any real commitment until its coefficients, loss scales, constraints, external outcomes, and calibration procedure have been pre-registered, externally validated, and independently replicated.

A full POS policy π over $[0, T]$ can be written as a constrained control problem:

$$\begin{aligned} J(\pi) = & E_{\pi}[Y_{\text{ext}}(T) + \lambda_U \cdot \int_{[0,T]} U \, dt - \lambda_B \cdot \int_{[0,T]} B \, dt - \lambda_C \cdot \int_{[0,T]} C_{\text{commit}} \, dt] \\ \text{s.t.} \quad & \text{Firewall}(a_t) = \text{PASS} \\ & G_{\text{POS}}(a_t, t) = 1 \quad \text{for every non-emergency commitment} \\ & C_{\text{total}}(\pi) \leq C_{\text{max}} \quad (\text{equal-budget constraint}) \\ \pi^*_{\text{POS}} = & \operatorname{argmax}_{\{\pi \in \Pi_{\text{POS}}\}} J(\pi) \end{aligned}$$

$Y_{\text{ext}}(T)$ external domain outcome, U internal stability, B base damage, C_{commit} commitment burden. This does **not** establish POS is optimal; it *defines the control problem* that POS-P1 and future simulations must compare against ungated and alternative-gated policies.

Epistemic level: B2 / L2. A serious, testable commitment discipline with two pre-registered falsifiers; **scaffolding until POS-P1/P2 report**, and marked as such.

9. Equation & variable registry

Symbol	Meaning	Status / units	Note
F, P, A	scores of the three pillars	dimensionless, usually $[0,1]$	not the physical currencies Time / Space / Energy
$U = \sqrt[3]{(F \cdot P \cdot A)}$	geometric aggregate	dimensionless	a stability <i>score</i> , not a hard constraint
a_t	a specific intervention at time t	domain-specific	distinct from the Action pillar score A_t
$\theta_F(a, d, s), \theta_P(a, d, s)$	readiness thresholds (action·domain·stakes)	dimensionless	must be empirically calibrated (POS-P1)
$\kappa(a)$	irreversibility / commitment degree	dimensionless, $[0,1]$	separates probes from commitments
$R(a, t)$	readiness ratio	dimensionless	readiness holds when $R \geq 1$
$G_{\text{POS}}(a, t)$	POS authorization gate	binary $\{0,1\}$	1 authorizes; 0 withholds
$B(a, t)$	base damage	normalized / domain-specific	requires a causal comparator (ΔB)
$\mathcal{M} = \int U \, dt$	accumulated stability integral	time-dimensioned unless normalized	cumulative stability / area under the U -trajectory
$\langle U \rangle_T = (1/T) \cdot \int U \, dt$	mean stability over $[0,T]$	dimensionless	the normalized companion to $\mathcal{M}$

Symbol	Meaning	Status / units	Note
$\mathcal{M}_{\text{net}}$	loss-adjusted stability stock	domain-specific	the <i>only</i> form that may decline (§9.1)
$L(a)$	expected loss of action a	domain-specific	same scale as $L(\text{inaction})$
w_F, w_P	base-damage weights	dimensionless	declared before analysis
m_F, m_P	normalized pillar readiness margins	dimensionless	positive above threshold, negative below
m_R	joint readiness margin	dimensionless	$m_R = R - 1$
$b(a, t)$	active readiness bottleneck	categorical (F or P)	identifies the limiting gate
$\alpha(s)$	stakes-conditioned confidence level	dimensionless (0,1)	decreases as stakes rise
$e(a)$	normalized action exposure	dimensionless [0,1]	separate from irreversibility
V_{POS}	readiness-violation severity	dimensionless / audit-scaled	does not replace the binary gate
M_{emg}	robust emergency margin	loss units	emergency support when $M_{\text{emg}} > 0$
$ER(a, t)$	emergency necessity ratio	dimensionless	robust support when $ER > 1$
c_F, c_P	costs of readiness verification	domain-specific	compare F-first vs P-first
p_F, p_P	readiness pass-probabilities	dimensionless [0,1]	estimated before ordering analysis
g_F, g_P	action-generated pillar gains	score change / step	constructive contribution
d_F, d_P	action-generated pillar damage	score change / step	destructive contribution
B_H	persistent base burn over horizon H	domain-specific / normalized	temporary vs durable damage
ΔM_T	counterfactual accumulation gap	time-dimensioned unless normalized	secondary internal outcome
$\Omega(a, t)$	option-foreclosure cost	value / loss units	model-dependent
C_{commit}	commitment burden	common decision units	direct cost + lost option value
$J(\pi)$	policy-level objective	domain utility	research formalization, not a validated law

9.1 The monotonicity of $\mathcal{M}$ (stated precisely)

For $U(t) \in [0,1]$ and $\mathcal{M}(t) = \int[0,t] U \, d\tau$, we have $d\mathcal{M}/dt = U(t) \geq 0$. $\mathcal{M}$ **cannot literally reverse or decrease** — it can only (i) grow more slowly, (ii) stall when $U = 0$, or (iii) fall behind a counterfactual that would have accumulated faster. So a "burn" harms the **growth** of $\mathcal{M}$, never its stock. Where the model genuinely needs a stock that can *fall*, use the explicit loss-adjusted form:

$$\mathcal{M}_{\text{net}}(t) = \mathcal{M}_0 + \int[0,t] (U(\tau) - L(\tau)) \, d\tau, \quad L(\tau) \geq 0$$

which declines exactly when $L(\tau) > U(\tau)$. Unqualified $\mathcal{M} = \int U \, d\tau$ is therefore always "cumulative stability / area under the U -trajectory"; any claim of a *declining* stock must name $\mathcal{M}_{\text{net}}$.

Epistemic status of the extended formal core (§§2.6.1–2.6.5, 3.3–3.4, 4.1–4.2, POS-P3).

These equations convert POS into a more explicit measurement-and-control specification. They do **not** upgrade its empirical status. The readiness margin, risk-adjusted thresholds, violation score, emergency margin, persistent-burn metric, counterfactual gap, option-foreclosure cost, and policy objective are all **L2/B2 candidate operationalizations**: their functional forms, coefficients, confidence rules, horizons, and loss scales require preregistration, calibration, ablation, and comparison against rival policies. *A formula makes a claim inspectable; it does not make the claim true.*

10. Integration into mathematical engines — POS as a plug-in commitment gate

POS is designed to drop into an existing engine as a **filter layer between candidate selection and irreversible execution** — it changes nothing in the discovery / optimization loop; it only decides whether an already-chosen action may be *committed* now. A runnable reference (pure-Python, standard library only, so it ports to C++/Julia/Rust or symbolic solvers) ships with this appendix: `pos_reference.py` (the gate, readiness metrics, base-burn, verification economics, reward-shaping helper) + `test_pos.py` (property tests).

The contract. *Inputs:* two readiness estimates with uncertainty (F , P as mean + standard error), and the candidate action's κ , e , s , domain, values-firewall verdict, and — for emergencies — loss bounds. *Output:* `authorize` $\in \{\text{true}, \text{false}\}$, a `reason`, and diagnostics (R , m_R , `bottleneck`, V_{POS} , M_{emg}). The gate is the faithful §2.6 predicate — **a high- κ commit is authorized only when $R \geq 1$** ; a reversible probe / low- κ prep is ungated; a values-forbidden action is refused at every readiness.

Where it plugs in, per engine:

Engine	POS insertion
Reinforcement learning	an action mask on irreversible commits ($G_{\text{POS}} = 0 \rightarrow \text{masked}$), plus reward shaping $r' = r + \lambda_U \cdot U - \lambda_B \cdot B - \lambda_V \cdot V_{\text{POS}}$; probes / prep stay available, so exploration is preserved
Model-predictive / optimal control	$G_{\text{POS}}(a_t, t) = 1$ as a hard constraint in the receding-horizon program; the POS-P3 objective $J(\pi)$ (§8) as the cost
Hierarchical planning	the gate between the planner (discovery + selection) and the executor — the plan is generated freely; the commit step must clear
Multi-agent systems	the Σ-agent commitment veto (§5): the F/P/A agents diagnose in parallel; Σ withholds <code>execute(a^commit)</code> until $G_{\text{POS}} = 1$
Discrete-event simulation	evaluate the gate at each commit event; log V_{POS} and every <code>a^emg</code> forced move for audit

Substrate-mapping recipe (what makes it interdisciplinary). (1) Map the domain's quantities to F , P , $A \in [0, 1]$ — normalized readiness of capability/structure, environment-fit/standing, operational readiness. (2) Attach κ , e , s to each action. (3) Calibrate $\theta_i(a, d, s)$ for the domain (POS-P1); declare $\alpha(s)$, the uncertainty model, and the loss scale **before** running the gate. (4) Wire the gate at the commit boundary. The *same* code then serves a trading system, a clinical protocol, a launch campaign, a governance reform, or an autonomous agent — only the F/P/A definitions and the θ calibration change.

A note on the naïve gate. The tempting shortcut "*if the action is decisive/irreversible, just authorize it*" is exactly inverted: high κ is what **raises** the readiness bar, never what waives it. The reference

implementation refuses a high- κ commit unless $R \geq 1$, precisely to avoid that error.

Scope (unchanged). §10 makes POS *runnable and embeddable*, not *validated*. The gate remains an L2/B2 candidate operationalization; a live deployment must first calibrate θ , α , κ , and the loss scales for its domain and clear POS-P1/P2. Embedding a formula is not evidence that the formula pays.

Reference code: `pos_reference.py`, `test_pos.py` (MIT; ship with this appendix).

References & provenance

Implementation companion. APPENDIX_POS_ENGINE_PROFILE defines the universal kernel, the domain-adapter contract, the uncertainty modes, the irreversibility estimators, the typed interfaces, and the compliance requirements. `pos_reference.py` (v0.2) is the RUO scalar reference implementation; `pos_example_lbo.py` is a **computational consistency illustration, not empirical evidence**.

- **Parent record:** U-Theory / U-Model — DOI [10.17605/OSF.IO/74XGR](https://doi.org/10.17605/OSF.IO/74XGR) · <https://u-model.org> · <https://github.com/UniversalModel>
- **Inherited primitives:** SSS (score U , δ , SI + stakes-scaled thresholds), GSI-RTD (AD-RTD $A \rightarrow F \rightarrow P$ candidate generation + the Triadic Scheduler, fed by π^* weakest-pillar triage + SSS-Guard), LGP (the 12-step cycle: parallel F/P/A diagnostics, LGP-7 selection, LGP-10 abort), TAA (the F/P/A/ Σ jury), TE (umbrella).
- **Sibling cross-refs:** MMT / DIM (the L3 meaning-interpretation layer; non-load-bearing here), QMC (measurement/irreversibility discussions), WAR / TRC / TRA firewalls (values-forbidden action), domain instances TRA / TRB / TRC / TSE (the TRA hardening supplies the zero-limit-only scoping of non-compensation adopted in §3).
- **Author:** Petar Nikolov (ORCID 0009-0001-8669-2276). © 2026, CC BY 4.0.

Version 1.6.1 — doc pass: F-first re-scoped as the canonical default (adapter may select P-first on pre-registered cost/dependency; §4), POS-P3 marked NOT-FOR-OPERATIONAL-AUTHORIZATION, and an implementation-companion cross-ref to APPENDIX_POS_ENGINE_PROFILE / `pos_reference.py` (v0.2). — v1.6 — POS made a **runnable, embeddable plug-in commitment gate** (§10): engine-integration contract + per-engine insertion patterns (RL action-mask & reward shaping, MPC hard constraint, hierarchical-planning gate, multi-agent Σ -veto, discrete-event) + a substrate-mapping recipe; ships a pure-Python reference `pos_reference.py` (faithful §2.6 predicate — high- κ commits require $R \geq 1$; Acklam inverse-normal so $\alpha(s)$ needs no SciPy) and property tests `test_pos.py` (13/13). — v1.5 **maximum formalization for interdisciplinary math engines** (10 compatible patches; binary gate & B2/L2 unchanged): readiness margin + active bottleneck (§2.6.1), explicit lower-confidence-bound standard with stakes-scaled α (§2.6.2), risk-adjusted threshold family $\theta_i(a, d, s)$ (§2.6.3), readiness-violation severity V_{POS} (§2.6.4), uncertainty-adjusted robust emergency override M_{emg}/ER (§2.6.5), base dynamics + persistent H-horizon burn B_H (§3.3), counterfactual accumulation gap ΔM_T (§3.4), F-first verification-cost economics (§4.1), commitment burden + option-foreclosure Ω (§4.2), the POS-P3 policy-level control objective $J(\pi)$ (§8), a full registry expansion (§9), and an epistemic guardrail keeping all of it L2/B2 candidate operationalizations. — v1.4 added the **formal core**: the gate as a computable predicate — readiness ratio $R(a, t)$ on lower confidence bounds, irreversibility $\kappa(a)$, the authorization gate $G_{POS}(a, t)$ with its regime table (§2.6); the base-burning metric $B(a, t)$ and its causal comparator ΔB (§3.2); and a full equation & variable registry with the precise $\mathcal{M}$ monotonicity

treatment and the loss-adjusted $\mathcal{M}_{\text{net}}$ (§9). — v1.3 minor revision per second re-review (ACCEPT, 8.6/10): F-first re-scoped as the default verification order (the $F\wedge P$ conjunction is order-independent; superiority is part of POS-P1) §4; a^{prep} gating made κ -based, not label-based, with an explicit no-label-gaming rule §2.5; $\mathcal{M}$ introduced as cumulative stability with the meaning-reading demoted to the L3 MMT/DIM layer §0/§6. — v1.2 major revision per adversarial re-review: POS repositioned as a **readiness-ordered commitment gate** (not a thinking-order law); reconciled with GSI-RTD's canonical AD-RTD $A \rightarrow F \rightarrow P$ discovery order (three-orders table §0, pipeline §5) and with LGP's parallel diagnostics (gate sits between LGP-7 and LGP-8/9); action taxonomy $a^{\text{probe}} / a^{\text{prep}} / a^{\text{commit}} / a^{\text{emg}} +$ irreversibility $\kappa(a)$ dissolves the action-circularity; energy wording corrected to commitment-irreversibility / option-value; $\mathcal{M}$ monotonicity corrected (growth collapses; the integral never reverses); non-compensation scoped to the zero-limit (the gate is a normative constraint, not a theorem of U); thresholds re-cast as $\vartheta(a, d, s)$ functions with LCBs, illustrative numbers stripped of decision authority; POS-P1/P2 upgraded to PASS / INCONCLUSIVE / KILL bands with mandatory AD-RTD baseline, ablations, external outcomes, and composite equal-budget; cosmological paragraph demoted to an L3 non-load-bearing note; license corrected. B2/L2 throughout. Not a guarantee of success — a discipline for accumulating stability (and, in the L3 MMT/DIM reading, meaning) over time.